

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
data=pd.read_csv("pd_speech_features.csv")
```

```
data.head()
```

↻

	id	gender	PPE	DFA	RPDE	numPulses	numPeriodsPulses	meanPeriodPulses	stdDevPeriodPulses	locPctJitter	...	tqwt_kurto
0	0	1	0.85247	0.71826	0.57227	240	239	0.008064	0.000087	0.00218	...	
1	0	1	0.76686	0.69481	0.53966	234	233	0.008258	0.000073	0.00195	...	
2	0	1	0.85083	0.67604	0.58982	232	231	0.008340	0.000060	0.00176	...	
3	1	0	0.41121	0.79672	0.59257	178	177	0.010858	0.000183	0.00419	...	
4	1	0	0.32790	0.79782	0.53028	236	235	0.008162	0.002669	0.00535	...	

5 rows × 755 columns

◀ ▶

```
data.corr() #Checking Correlation
```

↻

	id	gender	PPE	DFA	RPDE	numPulses	numPeriodsPulses	meanPeriodPulses	stdDevPeriodPulses	locPctJitter	...	tqwt_kurto
id	1.000000	-0.133605	0.026667	0.041938	-0.084606	0.085828	0.085226	-0.100831	-0.060000	-0.00218	...	-0.060000
gender	-0.133605	1.000000	0.010175	0.099356	0.168321	-0.478367	-0.477710	0.460422	-0.110000	0.00195	...	-0.110000
PPE	0.026667	0.010175	1.000000	-0.094775	-0.405558	0.191535	0.194098	-0.201907	-0.420000	0.00176	...	-0.420000
DFA	0.041938	0.099356	-0.094775	1.000000	0.155075	-0.286791	-0.286611	0.253708	0.080000	0.00419	...	0.080000
RPDE	-0.084606	0.168321	-0.405558	0.155075	1.000000	-0.521193	-0.524839	0.506707	0.330000	0.00535	...	0.330000
...
tqwt_kurtosisValue_dec_33	-0.053781	0.107618	0.045755	-0.033914	-0.097161	-0.047946	-0.047837	0.039487	-0.090000	0.00218	...	-0.090000
tqwt_kurtosisValue_dec_34	-0.058034	0.128936	0.046175	0.043454	-0.036406	-0.068664	-0.068561	0.046140	-0.060000	0.00195	...	-0.060000
tqwt_kurtosisValue_dec_35	-0.057807	0.107734	0.037385	0.067843	-0.018052	-0.069645	-0.069468	0.040828	-0.060000	0.00176	...	-0.060000
tqwt_kurtosisValue_dec_36	-0.055775	0.104828	0.039588	0.116699	0.021945	-0.062925	-0.062790	0.031685	-0.060000	0.00419	...	-0.060000
class	-0.111661	0.182713	-0.072939	0.306070	0.247444	-0.284056	-0.284002	0.211368	0.070000	0.00535	...	0.070000

755 rows × 755 columns

◀ ▶

```
data.describe() #Checking for Outliers
```

↻

	id	gender	PPE	DFA	RPDE	numPulses	numPeriodsPulses	meanPeriodPulses	stdDevPeriodPulses	locPctJitter	...	tqwt_kurto
count	756.000000	756.000000	756.000000	756.000000	756.000000	756.000000	756.000000	756.000000	756.000000	756.000000	756.000000	756.000000
mean	125.500000	0.515873	0.746284	0.700414	0.489058	323.972222	322.678571	0.006360	0.000383	0.00218	...	-0.060000
std	72.793721	0.500079	0.169294	0.069718	0.137442	99.219059	99.402499	0.001826	0.000728	0.00195	...	-0.110000
min	0.000000	0.000000	0.041551	0.543500	0.154300	2.000000	1.000000	0.002107	0.000011	0.00176	...	-0.420000
25%	62.750000	0.000000	0.762833	0.647053	0.386537	251.000000	250.000000	0.005003	0.000049	0.00419	...	0.080000
50%	125.500000	1.000000	0.809655	0.700525	0.484355	317.000000	316.000000	0.006048	0.000077	0.00535	...	0.330000
75%	188.250000	1.000000	0.834315	0.754985	0.586515	384.250000	383.250000	0.007528	0.000171	0.00535	...	0.330000
max	251.000000	1.000000	0.907660	0.852640	0.871230	907.000000	905.000000	0.012966	0.003483	0.00535	...	0.330000

8 rows × 755 columns

◀ ▶

```
X=data.drop("class",axis=1) #Splited the data in Dependent
y=data["class"] #Splited the data in Independent
```

```
!pip install xgboost
!pip install catboost
```

```
Requirement already satisfied: xgboost in /usr/local/lib/python3.10/dist-packages (2.1.3)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from xgboost) (1.26.4)
Requirement already satisfied: nvidia-nccl-cu12 in /usr/local/lib/python3.10/dist-packages (from xgboost) (2.23.4)
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from xgboost) (1.13.1)
Requirement already satisfied: catboost in /usr/local/lib/python3.10/dist-packages (1.2.7)
Requirement already satisfied: graphviz in /usr/local/lib/python3.10/dist-packages (from catboost) (0.20.3)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (from catboost) (3.8.0)
Requirement already satisfied: numpy<2.0,>=1.16.0 in /usr/local/lib/python3.10/dist-packages (from catboost) (1.26.4)
Requirement already satisfied: pandas>=0.24 in /usr/local/lib/python3.10/dist-packages (from catboost) (2.2.2)
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from catboost) (1.13.1)
Requirement already satisfied: plotly in /usr/local/lib/python3.10/dist-packages (from catboost) (5.24.1)
Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages (from catboost) (1.17.0)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.24->catboost) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.24->catboost) (2024.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.24->catboost) (2024.2)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (1.3.1)
Requirement already satisfied: cyclar>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (4.55.2)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (1.4.7)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (24.2)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (11.0.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (3.2.0)
Requirement already satisfied: tenacity>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from plotly->catboost) (9.0.0)
```

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=43) #Splited the data in test and train as 80% train and 20% t
```

```
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import accuracy_score
```

```
models={
    "RandomForestClassifier":RandomForestClassifier(),
    "XGBClassifier":XGBClassifier()
} #Checking with both models to know which can get more accuracy and precision
```

```
for i in models:
    model=models[i]
    model.fit(X_train,y_train)
    y_pred_test=model.predict(X_test)
    print(i)
    print("Accuracy_score:- ",accuracy_score(y_pred_test,y_test))
```

```
RandomForestClassifier
Accuracy_score:- 0.8448844884488449
XGBClassifier
Accuracy_score:- 0.8613861386138614
```

```
params={
    "n_estimators" : [110,100,150],
    "criterion" : ["gini","entropy"],
    "max_depth" : [1,2,3,4],
    "min_samples_leaf" : [1,2,3]
} # Custom Parameters Testing so that we can find best parameters for our models
rfc_model=XGBClassifier()
gd=GridSearchCV(rfc_model,params,cv=5,n_jobs=-1,verbose=2)
gd.fit(X_train,y_train)
```

```
Fitting 5 folds for each of 72 candidates, totalling 360 fits
/usr/local/lib/python3.10/dist-packages/xgboost/core.py:158: UserWarning: [07:52:27] WARNING: /workspace/src/learner.cc:740:
Parameters: { "criterion", "min_samples_leaf" } are not used.
```

```
warnings.warn(smsg, UserWarning)
```

```
GridSearchCV ⓘ ?
  ▸ best_estimator_: XGBClassifier
    ▸ XGBClassifier
```

```
gd.best_params_ # Best Parameters
```

```
{'criterion': 'gini',  
 'max_depth': 2,  
 'min_samples_leaf': 1,  
 'n_estimators': 100}
```

```
params={  
    "n_estimators" : 100,  
    "max_depth" : 3,  
    "min_samples_leaf" : 2  
}  
xgb_model=XGBClassifier(**params)  
xgb_model.fit(X_train,y_train)  
y_pred_test=xgb_model.predict(X_test)  
print("Accuracy_score:- ",accuracy_score(y_pred_test,y_test))
```

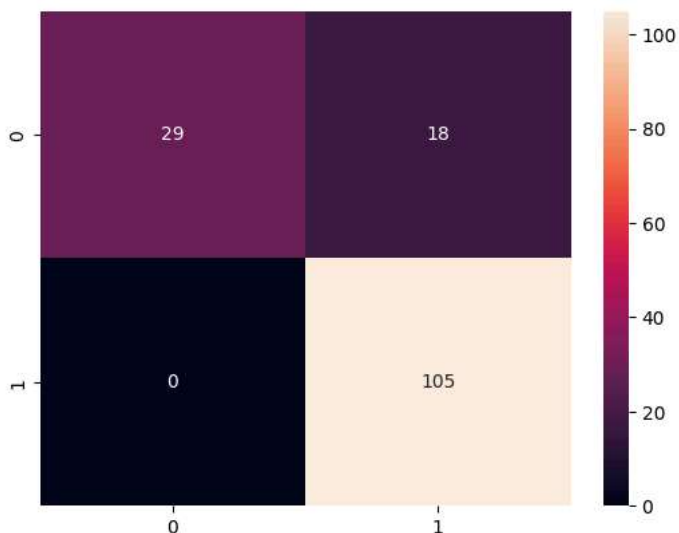
```
/usr/local/lib/python3.10/dist-packages/xgboost/core.py:158: UserWarning: [07:58:30] WARNING: /workspace/src/learner.cc:740:  
Parameters: { "min_samples_leaf" } are not used.
```

```
warnings.warn(smsg, UserWarning)  
Accuracy_score:- 0.881578947368421
```

```
from sklearn.metrics import classification_report,confusion_matrix  
print(classification_report(y_test,y_pred_test))  
sns.heatmap(confusion_matrix(y_test,y_pred_test),annot=True,fmt="d")
```

```
precision    recall  f1-score   support  
  
0           1.00      0.62      0.76         47  
1           0.85      1.00      0.92        105  
  
accuracy          0.88         152  
macro avg          0.93      0.81      0.84         152  
weighted avg          0.90      0.88      0.87         152
```


<Axes: >



```
from sklearn.metrics import accuracy_score  
print(accuracy_score(y_test,y_pred_test))
```

```
0.881578947368421
```

```
from joblib import load,dump  
dump(xgb_model,"xgb_model.joblib")
```

 ['xgb_model.joblib']